

DOMONDON DOMINIUM™

The MVV Field Guide

One team. One month. Twelve documents.

12

The Minimum Viable Version of a governed, AI-coordinated team —
no platform, no budget line, no permission needed.
Documents and resolve.

Carry the lamp. Keep the ledger. Agents propose. Humans judge. Nurses steward.

Robert Domondon, RN, CCRN · domondon-dominium.org

Before You Build

Read this page first. The month works because of four rules — and one act of honesty.

Take the vital signs first.

Twenty minutes, eight questions: how work arrives, who decides, what lives only in someone's head, what happens when things break. Do it on paper or at domondon-dominium.org/vitals.html. The verbs that hurt tell you which week below deserves extra care. The flinches are the map.

Rule 1 — Tell the truth.

The MVV is your team, finally described truthfully — folklore, workarounds, and all. A flattering charter is worthless. Week one is honesty; that was always the hard part.

Rule 2 — Documents, not tools.

Nothing here requires software. Every artifact is a page someone could print. If you later add AI coordination, these pages are exactly what the AI loads — a system can only be governed if it can describe itself.

Rule 3 — Every chain ends in a named human.

AI agents may draft and coordinate; they never own, approve, or review-of-record. Tasks delegate. Judgment does not. Accountability never moves.

Rule 4 — Schedule the loop before the month ends.

The twelfth artifact is a standing monthly appointment. A system that runs its improvement loop compounds; one that skips it accumulates silent risk with excellent documentation.

How to use the next four pages.

One week per page, three artifacts per week, a definition of done for each. Check the boxes as you go — or track online with your browser at domondon-dominium.org/mvv.html. Expect each artifact to take one honest hour.

Carry the lamp. Keep the ledger. Agents propose. Humans judge. Nurses steward.

Week 1 — Purpose & Promise

Why the team exists, whom it serves, what it offers. Judgment needs an object before it needs a tool.

1 Team charter

Who you are, why the team exists, and the intent clause — the two or three sentences that let a person or an AI make a sane call when the rules run out.

- ☐ Mission in one or two sentences
- ☐ Intent clause: the outcome + the non-negotiables
- ☐ Scope (what we own) and boundaries (what we never touch) listed

Tip: if your intent clause could belong to any team in the building, it isn't one.

2 Service catalog

Three services, each a published promise carrying its governance from birth.

- ☐ Three services named, each deliverable in a sentence
- ☐ Each carries: default risk tier · automation ceiling · evidence level · review gate
- ☐ The team agrees it can actually deliver all three

Tip: a catalog you can't deliver is a standing integrity failure — three is plenty.

3 Intake form / front door

One declared way in. Ambient demand — the hallway ask, the meeting aside — is where work disappears.

- ☐ One channel declared the front door; other channels route into it
- ☐ Form captures: what, for whom, by when, category
- ☐ The hallway ask has a rule: capture it always, promise it never

Tip: say yes to capturing the request — not to doing it.

Week 2 — The Work

How work is classified, executed, and owned. This is where governance meets the day shift.

4 Request triage

The governance checkpoint of first contact. A request's risk profile is set here and travels with it for life.

- ☐ Classification questions written: type, urgency, risk, owner, evidence, approval
- ☐ Governance profile fields: tier · data sensitivity · external-facing · human review
- ☐ P0–P4 plus Reject adopted — with the protection override in writing

Tip: anything touching safety, license, or protected data is P0 regardless of what scoring says.

5 Workflow map — one workflow

One real workflow made explicit, trigger to follow-up, with the automation ceiling declared per step.

- ☐ Steps written from trigger to documentation
- ☐ Each step declares: who · automation level (L0–L4) · evidence produced
- ☐ Review and approval are separate steps with different names on them

Tip: start with the workflow that embarrassed you most recently.

6 Role / owner map

The mixed human-AI roster. Eight role types; agents may hold exactly two of them.

- ☐ The eight role types assigned to actual names
- ☐ AI agents listed as Operator or Contributor only, each with a human principal
- ☐ Escalation owner named — and reachable

Tip: every chain of work terminates in a named human. Run your org chart against that sentence.

Carry the lamp. Keep the ledger. Agents propose. Humans judge. Nurses steward.

Week 3 — Memory & Control

What the team knows, and the gates that keep fluent output honest.

7 Knowledge source map

What the team runs on, tiered by authority: governed > working > ambient. Staleness gets an owner.

- ☐ Sources listed and tiered
- ☐ Each governed source has a named steward and a review date
- ☐ The things that live only in someone's head have a capture plan

Tip: an outdated policy served confidently is a latent failure. Date everything.

8 Human review gates

Not a human glancing at AI output — a designed control built not to rust open.

- ☐ Gates placed by risk and reversibility (and ungated-on-purpose written down)
- ☐ Each gate names its active task — what the reviewer does, not just reads
- ☐ Rejection counting started: a gate that never rejects gets re-examined

Tip: count the clicks. If approval is one click and rejection is a form, fix the form first.

9 The ledger

Decisions logged with rationale; actions attributable to named identities, human or agent.

- ☐ Append-only log created with an entry format
- ☐ First real entries written — including one gate event
- ☐ Rule adopted: no patient or confidential data in the ledger, ever

Tip: if it isn't in the ledger, it didn't happen — and can't teach anything.

Carry the lamp. Keep the ledger. Agents propose. Humans judge. Nurses steward.

Week 4 — The Heartbeat

Visibility, rhythm, and the loop that makes the whole system compound.

10 Task / status dashboard

Work visible across levels, carrying its governance profile with it.

- ☐ Active work listed with owner, status, due date
- ☐ Governance visible: "what high-risk work is currently ungated?" is answerable
- ☐ Stale items surface weekly, not archaeologically

Tip: a whiteboard counts. Visibility is the feature; software is optional.

11 Weekly brief — SBAR-shaped

Situation, Background, Assessment, Recommendation. A brief you can read in the elevator.

- ☐ Template set: S · B · A · R
- ☐ First brief produced and actually read
- ☐ A standing challenge round: disagreeing with the AI's brief is logged and praised

Tip: if the brief takes longer to write than the meeting it replaces, shorten the brief.

12 Metrics snapshot + improvement loop

The baseline — including governance health — and the monthly loop, on the calendar.

- ☐ Baseline taken: volume, speed, quality — whatever you honestly have
- ☐ Governance health included: gate events, escalations, boundary blocks
- ☐ First monthly improvement loop scheduled before the month ends

Tip: the loop reviews four things — workflows, prompts, knowledge, and the rules themselves.

After the Build

Twelve documents later, here is what to expect — and what to protect.

What changes first.

The team feels calmer and more honest. Work stops disappearing after meetings. Escalations stop being punished. The AI becomes useful in the boring, compounding way instead of the impressive, brittle way. People from other teams start asking what you did.

Watch one number.

Gate health. If your reviewers approve everything untouched for a month, the gate has rusted open — redesign the gate, not the reviewer. A reviewer who never rejects is a rubber stamp, and it is never her fault.

Keep one appointment.

The monthly improvement loop. One hour: one finding each for workflows, prompts, knowledge, and the governance rules. Install every lesson as a changed checklist, field, or permission — a fix that begins "remember to" is a wish with a calendar invite.

Go deeper.

The full framework — eighteen components, the Five Rights of AI Delegation, the autonomy ladder, the 24-item canvas with its maturity rubric — is free at domondon-dominium.org. The one-page Five Rights is at domondon-dominium.org/assets/five-rights.pdf. Print it. Pin it where your team delegates to machines.

Tasks delegate. Judgment does not. Accountability never moves.

Start with one team. Tell the truth about how it works. That's week one — and it was always the hard part.

Carry the lamp. Keep the ledger. Agents propose. Humans judge. Nurses steward.